

ANISH KAMATAM

📍 San Francisco Bay Area 📞 (925) 549-0603 📩 anishkamatam22@gmail.com 💬 linkedin.com/in/anishkamatam

Education

California State University - Chico <i>B.S. in Computer Science and Mathematics</i>	Expected May 2027
	Chico, CA

- **Coursework:** Operating Systems, Computer Networking, Theory of Computation, Data Structures & Algorithms, Robotics & Machine Intelligence, Applied Machine Learning, Artificial Intelligence, Applied Linear Algebra, Programming & Algorithms I & II

Skills

Python, C++, Go, Rust, Java, TypeScript, SQL, React, Node.js, FastAPI, PyTorch, TensorFlow, Scikit-Learn, TensorRT, Docker, Kubernetes, gRPC, Redis, PostgreSQL, AWS, GCP (Vertex AI), LangChain, LangGraph, Vector Databases

Experience

University of California, Berkeley – EECS <i>Undergraduate Research Assistant, Cloud Robotics Lab</i>	Oct 2025 – Present
	Berkeley, CA
<ul style="list-style-type: none">• Researching bandwidth-aware cloud robotics using ROS2, and PyTorch for offloading between Jetson Orin Nano and cloud GPUs.• Deployed transformer-based vision and SLAM models (DINOv2, YOLOv8) with TensorRT for low-latency perception over 5G.• Built Kubernetes-managed inference pipelines with adaptive compression, model switching, and telemetry profiling for drones.• Working under Prof. Sylvia Ratnasamy and Dr. Alex Krentsel on SLO-driven scheduling for scalable cloud autonomy.	
FoodHealth Company <i>AI/ML Engineer Intern</i>	Nov 2025 – Present
	San Francisco, CA
<ul style="list-style-type: none">• Trained XGBoost models on large tabular datasets, optimized hyperparameters with Optuna, reducing inference latency by 30%.• Built a quantile-scaled, target-encoded feature pipeline with polynomial interactions, improving model accuracy by 12%.• Developed an audit suite using SHAP, 5-fold CV, and KS-drift checks, increasing model reliability and drift detection by 25%.	
FundMore AI <i>AI Engineer Intern</i>	May 2025 – Oct 2025
	Toronto, Canada
<ul style="list-style-type: none">• Built a RAG AI assistant on Google Vertex AI using Sentence-BERT embeddings, Vector Search, and PostgreSQL for analytics.• Developed multi-tenant inference stack in TypeScript, Angular, and Node.js with Sequelize ORM and Redis caching for queries.• Implemented function calling and dynamic SQL generation for natural language access to real-time mortgage data.• Fine-tuned Gemini and PaLM models on financial datasets using Vertex AI Training and reinforcement-style feedback.• Designed agentic pipelines with Python orchestration and LLM subroutines for retrieval, summarization, and anomaly detection.	
AGNI X <i>Founder & Lead ML Engineer</i>	Apr 2025 – Sep 2025
	San Francisco, CA
<ul style="list-style-type: none">• Built an autonomous drone platform integrating thermal, RGB, and LiDAR sensing for wildfire detection and terrain mapping.• Deployed YOLOv8-S and thermal CNN heads on NVIDIA Jetson Orin Nano using ROS2 fusion for low-latency onboard inference.• Implemented visual-inertial EKF localization and Bayesian fusion for precise georeferencing and flame discrimination.• Developed an active learning loop with semi-supervised retraining and OTA updates to improve detection across environments.	
TraceCloud <i>Software Engineer Intern</i>	May 2024 – Sept 2024
	San Jose, CA
<ul style="list-style-type: none">• Developed a scalable full-stack QA automation framework using Selenium WebDriver, Node.js, and PostgreSQL.• Integrated REST APIs with continuous regression pipelines in Python and Jenkins, reducing test runtime by 40%.• Engineered a GraphRAG pipeline that modeled requirements, dependencies, and Jira artifacts as graph embeddings, enabling low-latency semantic retrieval and cutting manual trace analysis by 30%.	

Projects

ViewLingo Vapi, OpenAI, ElevenLabs, React Native, Mentra OS — YCombinator x Mentra Hackathon Winner	2025
<ul style="list-style-type: none">• Built a real-time AR translator using smart glasses with Whisper transcription and Vapi-powered conversational learning.• Developed an agentic memory pipeline that tracks user context and adapts next-day lessons based on learning history.	
SentraAI DeepL, Perplexity Sonar Pro, ElevenLabs, Apify, MCP, FastAPI — AWS MCP x A2A Hackathon Winner	2025
<ul style="list-style-type: none">• Developed a multilingual emergency response agent converting unstructured speech into structured MCP dispatch actions.• Integrated DeepL translation, Perplexity classification, and Apify retrieval for location-aware autonomous routing.	
MedWise Next.js, Flask, Neo4j, Pinecone, Hugging Face — SF Hacks 2x Category Winner	2025
<ul style="list-style-type: none">• Developed a graph-based RAG system linking 1,000+ drugs with embeddings in Pinecone for generics and interaction queries.• Integrated Neo4j graph traversal with Hugging Face rerankers to improve retrieval precision and clinical interpretability.	
Wizard JavaScript, Node.js, TypeScript, Docker, Postgres, Shell scripting, JSON parsing	2025
<ul style="list-style-type: none">• Built Wizard, a CLI debugging tool with sandboxed error reproduction, syscall tracing, and structured root-cause analysis.• Added automated fix workflows using semantic error classification, environment diffing, and safe command execution.	